

IERG 4330/ESTR 4316/IEMS 5730 Spring 2022

Homework 5

Release date: Apr 20, 2022

Due date: 11:59:00 pm, May 10, 2022

No late homework will be accepted!

Every Student **MUST** include the following statement, together with his/her signature in the submitted homework.

I declare that the assignment submitted on Elearning system is original except for source material explicitly acknowledged, and that the same or related material has not been previously submitted for another course. I also acknowledge that I am aware of University policy and regulations on honesty in academic work, and of the disciplinary guidelines and procedures applicable to breaches of such policy and regulations, as contained in the website

<http://www.cuhk.edu.hk/policy/academichonesty/>.

Signed (Student Lee) Date: 10-5-22

Name Chan Kai Yee SID 1153124983

Submission notice:

- Submit your homework via the elearning system

General homework policies:

A student may discuss the problems with others. However, the work a student turns in must be created **COMPLETELY** by oneself **ALONE**. A student may not share **ANY** written work or pictures, nor may one copy answers from any source other than one's own brain.

Each student **MUST LIST** on the homework paper the **name of every person he/she has discussed or worked with**. If the answer includes content from any other source, the student **MUST STATE THE SOURCE**. Failure to do so is cheating and will result in sanctions. Copying answers from someone else is cheating even if one lists their name(s) on the homework.

If there is information you need to solve a problem but the information is not stated in the problem, try to find the data somewhere. If you cannot find it, state what data you need, make a reasonable estimate of its value and justify any assumptions you make. You will be graded not only on whether your answer is correct but also on whether you have done an intelligent analysis.

Q1.

a)

Code:

```
1 from pyspark.sql import SQLContext
2 from pyspark import SparkConf, SparkContext
3 from graphframes import *
4 import pyspark.sql.functions as f
5
6 sc = SparkContext.getOrCreate()
7 sqlContext = SQLContext(sc)
8
9 actions = sqlContext.read.csv("hdfs:///user/s1155124983/hw5_q1/mooc_actions.tsv", header = True, sep = '\t')\
10 .withColumnRenamed("USERID", "src").withColumnRenamed("TARGETID", "dst")
11
12 mocc_vertices = sqlContext.read.csv("hdfs:///user/s1155124983/hw5_q1/vertices.tsv", header = True, sep = '\t')
13
14 mocc_g = GraphFrame(mocc_vertices, actions)
15
16 # mocc_g.inDegrees.show()
17 # mocc_g.vertices.show()
18
19 # Num. of Ver
20 num_v = mocc_g.vertices.count()
21 print(num_v)
22
23 # Num. of Users
24 num_u = mocc_g.vertices.filter("type = 'User'").count()
25 print(num_u)
26
27 # Num. of Course Activity
28 num_ca = mocc_g.vertices.filter("type = 'Course Activity'").count()
29 print(num_ca)
30
31 # the number of edges
32 num_e = mocc_g.edges.count()
33 print(num_e)
34
35 # the vertex with the Largest in-degree
36 mocc_g.inDegrees.orderBy(f.desc("inDegree")).limit(1).show()
37
38 # the vertex with the Largest out-degree
39 mocc_g.outDegrees.orderBy(f.desc("outDegree")).limit(1).show()
40
```

Output for (i) – (vi):

i: 7144

ii: 7047

iii: 97

iv: 411749

v and vi shown on the graph

```
7144
7047
97
411749
+-----+-----+
|  id|inDegree|
+-----+-----+
|7055|    19474|
+-----+-----+
```

```
+-----+-----+
|  id|outDegree|
+-----+-----+
|1181|      505|
+-----+-----+
```

b.)

Code:

```
1 from pyspark.sql import SQLContext
2 from pyspark import SparkConf, SparkContext
3 from graphframes import *
4 import pyspark.sql.functions as f
5
6 sc = SparkContext.getOrCreate()
7 sqlContext = SQLContext(sc)
8
9 actions = sqlContext.read.csv("hdfs:///user/s1155124983/hw5_q1/mocc_actions.tsv", header = True, sep = '\t')\
10 .withColumnRenamed("USERID", "src").withColumnRenamed("TARGETID", "dst")
11
12 mocc_vertices = sqlContext.read.csv("hdfs:///user/s1155124983/hw5_q1/vertices.tsv", header = True, sep = '\t')
13
14 mocc_g = GraphFrame(mocc_vertices, actions)
15
16 mocc_g_fil = mocc_g.filterEdges("TIMESTAMP >= 10000 and TIMESTAMP <= 50000").dropIsolatedVertices()
17
18 num_v = mocc_g_fil.vertices.count()
19 print(num_v)
20
21 num_e = mocc_g_fil.edges.count()
22 print(num_e)
23
```

Output:

Number of nodes = 49

Number of edges = 243

```
Log Type: stdout
Log Upload Time: Thu May 05 21:26:11 +0800 2022
Log Length: 7
49
243
```

c.)

Code:

```
1 from pyspark.sql import SQLContext
2 from pyspark import SparkConf, SparkContext
3 from graphframes import *
4 import pyspark.sql.functions as f
5
6 sc = SparkContext.getOrCreate()
7 sqlContext = SQLContext(sc)
8
9 actions = sqlContext.read.csv("hdfs:///user/s1155124983/hw5_q1/mooc_actions.tsv", header = True, sep = '\t')\
10 .withColumnRenamed("USERID", "src").withColumnRenamed("TARGETID", "dst")
11
12 mocc_vertices = sqlContext.read.csv("hdfs:///user/s1155124983/hw5_q1/vertices.tsv", header = True, sep = '\t')
13
14 mocc_g = GraphFrame(mocc_vertices, actions)
15
16 mocc_g_fil = mocc_g.filterEdges("TIMESTAMP >= 10000 and TIMESTAMP <= 50000").dropIsolatedVertices()
17
18 # i
19 print("i\n")
20 path_i = mocc_g_fil.find("(a)-[e1]->(b); (c) - [e2] -> (b)")\
21 .filter("a.id != c.id")\
22 .filter("e1.timestamp <= e2.timestamp")
23
24 path_i.show()
25
26 count_i = path_i.count()
27 print(count_i)
28
29 # ii
30 print("ii\n")
31 path_ii = mocc_g_fil.find("(a)-[e1]->(b); (b) - [e2] -> (c)")\
32 .filter("a.id != b.id and b.id != c.id")\
33 .filter("e1.timestamp <= e2.timestamp")
34
35 path_ii.show()
36
37 count_ii = path_ii.count()
38 print(count_ii)
39
40 # iii
41 print("iii\n")
42 path_iii = mocc_g_fil.find("(a)-[e1]->(c); (b) - [e4] -> (c); (a) - [e3] -> (d); (b) - [e2] -> (d)")\
43 .filter("a.id != b.id and c.id != d.id")\
44 .filter("e1.timestamp <= e2.timestamp and e2.timestamp <= e3.timestamp and e3.timestamp <= e4.timestamp")
45
46 path_iii.show()
47
48 path_iii = path_iii.count()
49 print(path_iii)
50
51 # iv
52 print("iv\n")
53 path_iv = mocc_g_fil.find("(d)-[e1]->(a); (b) - [e3] -> (c); (d) - [e4] -> (c); (d) - [e2] -> (e)")\
54 .filter("a.id != c.id and b.id != d.id and c.id != e.id and a.id != e.id")\
55 .filter("e1.timestamp <= e2.timestamp and e2.timestamp <= e3.timestamp and e3.timestamp <= e4.timestamp")
56
57 path_iv.show()
58
59 path_iv = path_iv.count()
60 print(path_iv)
61
```

i)	2372
ii)	0
iii)	215
iv)	7005

		a)		e1)		b)		c)		e2)	
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[33, user]	[33, 7048, 49902, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[33, user]	[33, 7048, 49904, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[32, user]	[32, 7048, 49898, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[31, user]	[31, 7048, 45788, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[30, user]	[30, 7048, 45890, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[29, user]	[29, 7048, 43620, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[29, user]	[29, 7048, 43615, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[28, user]	[28, 7048, 42459, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[27, user]	[27, 7048, 41947, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[27, user]	[27, 7048, 41934, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[26, user]	[26, 7048, 41882, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[26, user]	[26, 7048, 41894, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[25, user]	[25, 7048, 41049, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[24, user]	[24, 7048, 40670, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[23, user]	[23, 7048, 40235, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[21, user]	[21, 7048, 39627, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[20, user]	[20, 7048, 39606, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[14, user]	[14, 7048, 39447, 0]							
[2, user]	[2, 7048, 37868, 0]	[7048, Course Act...	[14, user]	[14, 7048, 39438, 0]							

2372
110
111

i	a)	e1	c	b)	e4	e3	d)	e2
[6, user]	[6, 7048, 38218.0]	[7048, Course Act...	[14, user]	[14, 7048, 39447.0]	[6, 7060, 39445.0]	[7060, Course Act...	[14, 7060, 39133.0]	
[6, user]	[6, 7050, 38261.0]	[7048, Course Act...	[7, user]	[7, 7050, 38583.0]	[6, 7048, 38380.0]	[7048, Course Act...	[7, 7048, 38207.0]	
[6, user]	[6, 7048, 38380.0]	[7048, Course Act...	[6, user]	[6, 7048, 38447.0]	[6, 7050, 38445.0]	[7050, Course Act...	[14, 7060, 39133.0]	
[6, user]	[6, 7048, 38327.0]	[7048, Course Act...	[14, user]	[14, 7048, 39513.0]	[7060, 39445.0]	[7060, Course Act...	[14, 7060, 39133.0]	
[6, user]	[6, 7054, 38340.0]	[7054, Course Act...	[13, user]	[13, 7054, 39591.0]	[6, 7060, 39445.0]	[7060, Course Act...	[13, 7060, 39105.0]	
[6, user]	[6, 7054, 38606.0]	[7054, Course Act...	[13, user]	[13, 7054, 39591.0]	[6, 7060, 39445.0]	[7060, Course Act...	[13, 7060, 39105.0]	
[6, user]	[6, 7056, 38623.0]	[7056, Course Act...	[13, user]	[13, 7056, 39627.0]	[6, 7060, 39445.0]	[7060, Course Act...	[13, 7060, 39105.0]	
[6, user]	[6, 7056, 38623.0]	[7056, Course Act...	[13, user]	[13, 7056, 39627.0]	[6, 7060, 39445.0]	[7060, Course Act...	[13, 7060, 39105.0]	
[14, user]	[14, 7049, 39065.0]	[7049, Course Act...	[16, user]	[16, 7049, 41611.0]	[14, 7048, 39447.0]	[7048, Course Act...	[16, 7048, 39193.0]	
[14, user]	[14, 7049, 39065.0]	[7049, Course Act...	[16, user]	[16, 7049, 41611.0]	[14, 7048, 39370.0]	[7048, Course Act...	[16, 7048, 39193.0]	
[13, user]	[13, 7051, 39105.0]	[7051, Course Act...	[14, user]	[14, 7051, 39627.0]	[13, 7056, 39623.0]	[7056, Course Act...	[14, 7056, 39157.0]	
[13, user]	[13, 7051, 39105.0]	[7051, Course Act...	[14, user]	[14, 7051, 39627.0]	[13, 7056, 39623.0]	[7056, Course Act...	[14, 7056, 39133.0]	
[13, user]	[13, 7051, 39105.0]	[7051, Course Act...	[14, user]	[14, 7051, 39627.0]	[13, 7054, 39591.0]	[7054, Course Act...	[14, 7054, 39162.0]	
[13, user]	[13, 7051, 39105.0]	[7051, Course Act...	[14, user]	[14, 7051, 39624.0]	[13, 7056, 39623.0]	[7056, Course Act...	[14, 7056, 39157.0]	
[13, user]	[13, 7051, 39105.0]	[7051, Course Act...	[14, user]	[14, 7051, 39624.0]	[13, 7056, 39623.0]	[7056, Course Act...	[14, 7056, 39157.0]	
[13, user]	[13, 7051, 39105.0]	[7051, Course Act...	[14, user]	[14, 7051, 39624.0]	[13, 7056, 39623.0]	[7056, Course Act...	[14, 7056, 39157.0]	
[14, user]	[14, 7050, 39113.0]	[7050, Course Act...	[17, user]	[17, 7050, 41319.0]	[14, 7048, 39447.0]	[7048, Course Act...	[17, 7048, 39220.0]	
[14, user]	[14, 7050, 39113.0]	[7050, Course Act...	[17, user]	[17, 7050, 41319.0]	[14, 7048, 39438.0]	[7048, Course Act...	[17, 7048, 39220.0]	

215
iv[illegible]

7005

Q2

a.)

Code:

```
1 import org.apache.spark._
2 import org.apache.spark.graphx._
3 import org.apache.spark.rdd.RDD
4 import org.apache.spark.SparkContext
5 import org.apache.spark.graphx.GraphLoader
6
7 object SimpleApp {
8
9   def max(a: (VertexId, Int), b: (VertexId, Int)): (VertexId, Int) = {
10     if (a._2 > b._2) a else b
11   }
12
13   def main(args: Array[String]) {
14
15     val sc = new SparkContext()
16
17     val cite_edge = GraphLoader.edgeListFile(sc, "edge_list.txt")
18
19     val num_vert = cite_edge.vertices.count()
20     println(num_vert)
21
22     val num_edges = cite_edge.edges.count()
23     println(num_edges)
24
25     val vert_lar_in_d = cite_edge.inDegrees.reduce(max)
26     println(vert_lar_in_d)
27
28     val vert_lar_ouy_d = cite_edge.outDegrees.reduce(max)
29     println(vert_lar_ouy_d)
30
31   }
32 }
33
34
```

Output:

The answers follow the order of variables:

num_vert = 169343

num_edges = 1166243

vert_lar_in_d = 1353

vert_lar_ouy_d = 72253

b.)

Code:

```
1 import org.apache.spark._
2 import org.apache.spark.graphx._
3 import org.apache.spark.rdd.RDD
4 import org.apache.spark.SparkContext
5 import org.apache.spark.graphx.GraphLoader
6 import org.apache.spark.graphx.lib.PageRank
7
8 object SimpleApp{
9
10   def max(a: (VertexId, Int), b: (VertexId, Int)): (VertexId, Int) = {
11     if (a._2 > b._2) a else b
12   }
13
14   def main(args: Array[String]) {
15
16     val sc = new SparkContext()
17
18     val cite_edge = GraphLoader.edgeListFile(sc, "edge_list.txt")
19
20     val num_vert = cite_edge.vertices.count()
21     println(num_vert)
22
23     val num_edges = cite_edge.edges.count()
24     println(num_edges)
25
26     val vert_lar_in_d = cite_edge.inDegrees.reduce(max)
27     println(vert_lar_in_d)
28
29     val vert_lar_ouy_d = cite_edge.outDegrees.reduce(max)
30     println(vert_lar_ouy_d)
31
32     val conn_vert = cite_edge.connectedComponents().vertices
33
34     val same_conn = conn_vert.map((v: (Long, Long)) => v._2).distinct.count()
35     val conn_num = conn_vert.distinct.count()
36     println("bi")
37     println(same_conn)
38     println(conn_num)
39
40     println("bii")
41     val st_conn_vert = cite_edge.stronglyConnectedComponents(3).vertices.map(_._2).distinct.count()
42     println(st_conn_vert)
43   }
```

i.)

Connected components exist in the citation network = 1

Vertices exist in the largest connected component = 169343

ii.)

Strongly connected components exist in the citation network = 169343

#In line 41 the stronglyConnectedComponents function is changed to 1 instead of 3 because of the driver-memory problem.

c)

Code:

```
44 println("ci")
45 val pr_4300 = PageRank.runParallelPersonalizedPageRank(cite_edge, 10, 0.15, Array(4300,5730))
46
47 pr_4300.vertices.top(20)(Ordering.by(_._2(0))).foreach(println)
48
49 pr_4300.vertices.top(20)(Ordering.by(_._2(1))).foreach(println)
```

Output:

i.)

4330:

```
(4300, (2, [0, 1], [0.21727664180027045, 0.0]))
(154342, (2, [0, 1], [0.02973885013916902, 0.0]))
(115359, (2, [0, 1], [0.02960590277825054, 0.0]))
(60030, (2, [0, 1], [0.02030554143357159, 0.0]))
(124635, (2, [0, 1], [0.01856063264049989, 0.0]))
(88323, (2, [0, 1], [0.018077909174157376, 0.0]))
(37909, (2, [0, 1], [0.01628208343442753, 0.0]))
(1580, (2, [0, 1], [0.015025749345463362, 0.0]))
(141153, (2, [0, 1], [0.014619915854653557, 0.0]))
(7805, (2, [0, 1], [0.014490170774386262, 0.0]))
(40166, (2, [0, 1], [0.014345084440027623, 0.0]))
(136616, (2, [0, 1], [0.014262428507930398, 0.0]))
(137062, (2, [0, 1], [0.01420916900617627, 0.0]))
(112716, (2, [0, 1], [0.014206551062774585, 0.0]))
(12084, (2, [0, 1], [0.014206549656476558, 0.0]))
(57425, (2, [0, 1], [0.014206549656171314, 0.0]))
(159030, (2, [0, 1], [0.013328970873926397, 0.0]))
(135057, (2, [0, 1], [0.01175699580118964, 0.0]))
(137083, (2, [0, 1], [0.01111078223355775, 0.0]))
(92833, (2, [0, 1], [0.010043016367521203, 0.0]))
```

5730:

```
scala> pr_4300.vertices.top(20)(Ordering.by(_._2(1))).foreach(println)
(5730, (2, [0, 1], [0.0, 0.4522328999434708]))
(102862, (2, [0, 1], [0.0, 0.1921989824759751]))
(165911, (2, [0, 1], [0.0, 0.1921989824759751]))
(141857, (2, [0, 1], [0.0, 0.16336913510457884]))
(108150, (2, [0, 1], [0.0, 0.0]))
(68522, (2, [0, 1], [0.0, 0.0]))
(91902, (2, [0, 1], [0.0, 0.0]))
(38926, (2, [0, 1], [0.0, 0.0]))
(139526, (2, [0, 1], [0.0, 0.0]))
(32676, (2, [0, 1], [3.715000307771759E-6, 0.0]))
(154038, (2, [0, 1], [0.0, 0.0]))
(51620, (2, [0, 1], [0.0, 0.0]))
(23776, (2, [0, 1], [0.0, 0.0]))
(129434, (2, [0, 1], [0.0, 0.0]))
(153030, (2, [0, 1], [0.0, 0.0]))
(53926, (2, [0, 1], [0.0, 0.0]))
(103184, (2, [0, 1], [0.0, 0.0]))
(4926, (2, [0, 1], [0.0, 0.0]))
(63852, (2, [0, 1], [0.0, 0.0]))
(161980, (2, [0, 1], [0.0, 0.0]))
```

ii)

Code:

```
52     println("cii")
53     val pr_5730 = pr_4300.vertices.top(2000)(Ordering.by(_._2(1)))
54     val sub_5730 = cite_edge.subgraph(vpred = (id, attr) => pr_5730.map(_._1) contains id)
55
56     val sub_5730_ed_count = sub_5730.edges.count()
57     println(sub_5730_ed_count)
58
```

number of edges = 155

d.)

Code:

```
58  
59     println("d")  
60     val lp_cite = LabelPropagation.run(cite_edge, 50)  
61  
62     val dist_label = lp_cite.vertices.map(_._2).distinct.count()  
63     println(dist_label);  
64  
65     val max_comm = lp_cite.vertices.map(_._2).map((_,1)).reduceByKey(_+_).reduce(max)  
66     println(max_comm)  
67  
68 }  
69 }
```

Number of communities = 14107

Vertices in the largest communities = 49058

```
14107  
d  
14107  
(69794,49058)
```

e.)

Code:

```
1 import org.apache.spark._
2 import org.apache.spark.graphx._
3 import org.apache.spark.rdd.RDD
4 import org.apache.spark.SparkContext
5 import org.apache.spark.graphx.GraphLoader
6 import org.apache.spark.graphx.lib.PageRank
7 import org.apache.spark.graphx.lib.LabelPropagation
8
9 object SimpleApp2{
10
11   def main(args: Array[String]) {
12
13     val sc = new SparkContext()
14
15     val dag = GraphLoader.edgeListFile(sc, "dag_edge_list.txt")
16     val init_g = dag.mapVertices((_,_) => 0)
17
18     val sssp = init_g.pregel(0)(
19       (id, dist, newDist) => math.max(dist, newDist), // Vertex Program
20       triplet => { // Send Message
21         if (triplet.srcAttr + 1 > triplet.dstAttr) {
22           Iterator((triplet.dstId, triplet.srcAttr + 1))
23         } else {
24           Iterator.empty
25         }
26       },
27       (a, b) => math.max(a, b) // Merge Message
28     )
29
30     println(sssp.vertices.collect.mkString("\n"))
31   }
32 }
33
34
35
```

Output:

ՍԻՍՏԵՄ ԶԱՅՄԱՆՈՒՄ (ԼԱՌԱՆՈՒՄ - Ը. ԴՊՈՒՄՆԵՐԸ Ե՝ ԵՐԵՐ Ե՝ ԼՅՂ.՝

Log Type: stdout

Log Upload Time: Sun May 08 02:35:29 +0800 2022

Log Length: 1751757

Showing 4096 bytes of 1751757 total. Click [here](#) for the full log.

1501, 4)
(105239, 3)
(133503, 5)
(101361, 6)
(157119, 1)
(98321, 3)
(166149, 4)
(155253, 3)
(563, 4)
(71689, 3)
(160335, 3)
(147889, 4)
(91334, 2)
(55915, 4)
(50485, 1)
(48161, 3)
(84319, 2)
(50477, 0)
(90297, 4)
(19947, 2)
(20917, 3)
(50249, 3)
(73665, 3)
(43439, 2)
(138393, 2)
(151363, 4)
(25827, 4)
(156003, 5)
(84535, 5)
(109735, 7)
(26813, 3)
(101807, 3)
(5829, 0)
(135283, 0)
(79247, 3)
(161731, 1)
(162611, 5)
(6791, 2)
(17229, 3)
(66355, 2)
(117285, 2)
(52759, 2)
(27673, 5)
(154149, 3)
(104575, 4)
(95785, 4)
(81807, 1)
(162009, 5)
(29649, 3)
(57181, 4)
(143455, 4)
(58465, 1)
(82139, 0)
(93763, 3)
(154151, 5)
(90231, 3)
(83163, 6)
(65041, 0)
(28819, 3)
(16161, 6)

Q3.

a)

Adding row key and zero padding for occurrence count :

```
1 import numpy as np
2 i = 1
3
4 with open('g2','a') as g2:
5     with open('googlebooks-eng-all-1gram-20120701-b', 'r') as f:
6         lines = f.readlines()
7
8         for line in lines:
9             line = line.strip().split('\t')
10
11             line[2] = line[2].zfill(5)
12             line.insert(0, int(i))
13             np.savetxt(g2, [line], delimiter='\t', fmt = '%s', newline = '\n')
14             i = i + 1
15
```

```
[s1155124983@dicvmd10 hw5]$ more g2
1      B'enard 1974      00001      1
2      B'enard 1982      00001      1
3      B'enard 1993      00001      1
4      B'enard 1997      00001      1
5      B'enard 2001      00002      1
6      B'enard 2003      00003      2
7      B'enard 2004      00008      6
8      B'enard 2005      00025      6
9      B'enard 2006      00156     14
10     B'enard 2007      00159     19
11     B'enard 2008      00044     18
12     B'h_NOUN      1794      00001      1
13     B'h_NOUN      1855      00001      1
14     B'h_NOUN      1872      00001      1
15     B'h_NOUN      1878      00003      2
16     B'h_NOUN      1884      00001      1
17     B'h_NOUN      1885      00001      1
18     B'h_NOUN      1892      00001      1
19     B'h_NOUN      1896      00001      1
```

ImportTsv:

```
[s1155124983@dicvmd10 hw5]$ hbase org.apache.hadoop.hbase.mapreduce.ImportTsv -Dimporttsv.columns=HBASE_ROW_KEY,cf:bigram,cf:year,cf:match_count,cf:vol_count -Dimporttsv.bulk.output=hdfs:///user/s1155124983/g3_fd g3_tb hdfs:///user/s1155124983/g2
```

Complete bulk load:

```
Bytes Written=11262865105
[s1155124983@dicvmd10 hw5]$ hbase org.apache.hadoop.hbase.mapreduce.LoadIncrementalHFiles hdfs:///user/s1155124983/g3_fd g3_tb
```

Showing the table:

```
hbase(main):034:0> scan 'g3 tb', {'LIMIT' => 20}
ROW      COLUMN+CELL
1        column=cf:bigram, timestamp=1652186559702, value=B'enard
1        column=cf:match_count, timestamp=1652186559702, value=00001
1        column=cf:vol_count, timestamp=1652186559702, value=1
1        column=cf:year, timestamp=1652186559702, value=1974
10       column=cf:bigram, timestamp=1652186559702, value=B'enard
10       column=cf:match_count, timestamp=1652186559702, value=00159
10       column=cf:vol_count, timestamp=1652186559702, value=19
10       column=cf:year, timestamp=1652186559702, value=2007
100      column=cf:bigram, timestamp=1652186559702, value=B.138 NOUN
100      column=cf:match_count, timestamp=1652186559702, value=00001
100      column=cf:vol_count, timestamp=1652186559702, value=1
100      column=cf:year, timestamp=1652186559702, value=1945
1000     column=cf:bigram, timestamp=1652186559702, value=Bismarck
1000     column=cf:match_count, timestamp=1652186559702, value=00002
1000     column=cf:vol_count, timestamp=1652186559702, value=2
1000     column=cf:year, timestamp=1652186559702, value=1876
10000    column=cf:bigram, timestamp=1652186559702, value=BON
10000    column=cf:match_count, timestamp=1652186559702, value=00066
10000    column=cf:vol_count, timestamp=1652186559702, value=58
10000    column=cf:year, timestamp=1652186559702, value=1863
100000   column=cf:bigram, timestamp=1652186559702, value=battery.2 NOUN
```

b.)

1. Insert ierg4330 2019 100 4

```
hbase(main):035:0> put 'g3_tb', '0', 'cf:bigram', 'ierg4330'
0 row(s) in 0.0100 seconds

hbase(main):036:0> put 'g3_tb', '0', 'cf:year', '2019'
0 row(s) in 0.0080 seconds

hbase(main):037:0> put 'g3_tb', '0', 'cf:match_count', '00100'
0 row(s) in 0.0310 seconds

hbase(main):038:0> put 'g3_tb', '0', 'cf:vol_count', '4'
0 row(s) in 0.0080 seconds
```

Updated table:

```
hbase(main):039:0> scan 'g3_tb', {'LIMIT' => 20}
ROW COLUMN+CELL
0 column=cf:bigram, timestamp=1652187371724, value=ierg4330
0 column=cf:match_count, timestamp=1652187407025, value=00100
0 column=cf:vol_count, timestamp=1652187416357, value=4
0 column=cf:year, timestamp=1652187396203, value=2019
1 column=cf:bigram, timestamp=1652186559702, value=B'enard
1 column=cf:match_count, timestamp=1652186559702, value=00001
1 column=cf:vol_count, timestamp=1652186559702, value=1
1 column=cf:year, timestamp=1652186559702, value=1974
10 column=cf:bigram, timestamp=1652186559702, value=B'enard
10 column=cf:match_count, timestamp=1652186559702, value=00159
```


2. Filter

```
hbase(main):012:0> scan 'g3 tb', {FILTER => "SingleColumnValueFilter('cf','year',=,'binary:1671') AND SingleColumnValueFilter('cf','match_count',>,'binary:00100')",COLUMNS => ['cf']}
ROW COLUMN+CELL
1839085 column=cf:bigram, timestamp=1652186559702, value=but
1839085 column=cf:match_count, timestamp=1652186559702, value=00644
1839085 column=cf:vol_count, timestamp=1652186559702, value=4
1839085 column=cf:year, timestamp=1652186559702, value=1671
20302314 column=cf:bigram, timestamp=1652186559702, value=but_CONJ
20302314 column=cf:match_count, timestamp=1652186559702, value=00643
20302314 column=cf:vol_count, timestamp=1652186559702, value=4
20302314 column=cf:year, timestamp=1652186559702, value=1671
23178288 column=cf:bigram, timestamp=1652186559702, value=been
23178288 column=cf:match_count, timestamp=1652186559702, value=00237
23178288 column=cf:vol_count, timestamp=1652186559702, value=4
23178288 column=cf:year, timestamp=1652186559702, value=1671
2401774 column=cf:bigram, timestamp=1652186559702, value=be
2401774 column=cf:match_count, timestamp=1652186559702, value=01230
2401774 column=cf:vol_count, timestamp=1652186559702, value=4
2401774 column=cf:year, timestamp=1652186559702, value=1671
29638627 column=cf:bigram, timestamp=1652186559702, value=being
```

3. Deleted the records in part 2

```
hbase(main):013:0> delete 'g3 tb', {FILTER => "SingleColumnValueFilter('cf','year',!=,'binary:1671') OR SingleColumnValueFilter('cf','match_count',<,'binary:00100')"}
1 row(s) deleted.
```

```
10007103 column=cf:vol_count, timestamp=1652186559702, value=1
10007103 column=cf:year, timestamp=1652186559702, value=1770
10007104 column=cf:bigram, timestamp=1652186559702, value=bylines_VERB
10007104 column=cf:match_count, timestamp=1652186559702, value=00001
10007104 column=cf:vol_count, timestamp=1652186559702, value=1
10007104 column=cf:year, timestamp=1652186559702, value=1829
10007105 column=cf:bigram, timestamp=1652186559702, value=bylines_VERB
10007105 column=cf:match_count, timestamp=1652186559702, value=00001
10007105 column=cf:vol_count, timestamp=1652186559702, value=1
10007105 column=cf:year, timestamp=1652186559702, value=1860
10007106 column=cf:bigram, timestamp=1652186559702, value=bylines_VERB
10007106 column=cf:match_count, timestamp=1652186559702, value=00001
10007106 column=cf:vol_count, timestamp=1652186559702, value=1
10007106 column=cf:year, timestamp=1652186559702, value=1875
10007107 column=cf:bigram, timestamp=1652186559702, value=bylines_VERB
10007107 column=cf:match_count, timestamp=1652186559702, value=00001
10007107 column=cf:vol_count, timestamp=1652186559702, value=1
10007107 column=cf:year, timestamp=1652186559702, value=1880
10007108 column=cf:bigram, timestamp=1652186559702, value=bylines_VERB
10007108 column=cf:match_count, timestamp=1652186559702, value=00001
10007108 column=cf:vol_count, timestamp=1652186559702, value=1
10007108 column=cf:year, timestamp=1652186559702, value=1887
10007109 column=cf:bigram, timestamp=1652186559702, value=bylines_VERB
10007109 column=cf:match_count, timestamp=1652186559702, value=00004
10007109 column=cf:vol_count, timestamp=1652186559702, value=4
10007109 column=cf:year, timestamp=1652186559702, value=1893
1000711 column=cf:bigram, timestamp=1652186559702, value=Bromley.1
1000711 column=cf:match_count, timestamp=1652186559702, value=00003
```